

```
1 #%% md
2 # Intelligent Agents: Reflex-Based Agents for the
3 # Vacuum-cleaner World
4
5 ## Instructions
6
7 Total Points: Undergrads 100 / Graduate students
8 110
9 Complete this notebook. Use the provided notebook
10 cells and insert additional code and markdown cells
11 as needed. Submit the completely rendered notebook
12 as a PDF file.
13 In this assignment you will implement a simulator
14 environment for an automatic vacuum cleaner robot,
15 a set of different reflex-based agent programs, and
16 perform a comparison study for cleaning a single
17 room. Focus on the __cleaning phase__ which starts
18 when the robot is activated and ends when the last
19 dirty square in the room has been cleaned. Someone
20 else will take care of the agent program needed to
21 navigate back to the charging station after the
22 room is clean.
23
24 ## PEAS description of the cleaning phase
25
26 __Performance Measure:__ Each action costs 1 energy
27 unit. The performance is measured as the sum of
28 the energy units used to clean the whole room.
29
30 __Environment:__ A room with  $n \times n$  squares
31 where  $n = 5$ . Dirt is randomly placed on each
32 square with probability  $p = 0.2$ . For simplicity,
33 you can assume that the agent knows the size and
34 the layout of the room (i.e., it knows  $n$ ). To
35 start, the agent is placed on a random square.
```

```

21 __Actuators__: The agent can clean the current
    square (action `suck`) or move to an adjacent
    square by going `north`, `east`, `south`, or `west`
    .
22
23 __Sensors__: Four bumper sensors, one for north,
    east, south, and west; a dirt sensor reporting dirt
    in the current square.
24
25
26 ## The agent program for a simple randomized agent
27
28 The agent program is a function that gets sensor
    information (the current percepts) as the arguments
    . The arguments are:
29
30 * A dictionary with boolean entries for the for
    bumper sensors `north`, `east`, `west`, `south`. E.
    g., if the agent is on the north-west corner,
    bumpers` will be `{"north" : True, "east" : False
    , "south" : False, "west" : True}`.
31 * The dirt sensor produces a boolean.
32
33 The agent returns the chosen action as a string.
34
35 Here is an example implementation for the agent
    program of a simple randomized agent:
36 #%%
37 import numpy as np
38
39 actions = ["north", "east", "west", "south", "suck"]
40
41 def simple_randomized_agent(bumpers, dirty):
42     return np.random.choice(actions)
43 #%%
44 # define percepts (current location is NW corner
        and it is dirty)
45 bumpers = {"north" : True, "east" : False, "south"
            : False, "west" : True}
46 dirty = True

```

```

47
48 # call agent program function with percepts and it
   returns an action
49 simple_randomized_agent(bumpers, dirty)
50 #%% md
51 __Note:__ This is not a rational intelligent agent
   . It ignores its sensors and may bump into a wall
   repeatedly or not clean a dirty square. You will be
   asked to implement rational agents below.
52 #%% md
53 ## Simple environment example
54
55 We implement a simple simulation environment that
   supplies the agent with its percepts.
56 The simple environment is infinite in size (bumpers
   are always `False`) and every square is always
   dirty, even if the agent cleans it. The environment
   function returns a performance measure which is
   here the number of cleaned squares (since the room
   is infinite and all squares are constantly dirty,
   the agent can never clean the whole room as
   required in the PEAS description above). The energy
   budget of the agent is specified as `max_steps`.
57 #%%
58 def simple_environment(agent, max_steps, verbose =
   True):
59     num_cleaned = 0
60
61     for i in range(max_steps):
62         dirty = True
63         bumpers = {"north" : False, "south" : False
   , "west" : False, "east" : False}
64
65         action = agent(bumpers, dirty)
66         if (verbose): print("step", i , "- action:"
   , action)
67
68         if (action == "suck"):
69             num_cleaned = num_cleaned + 1
70
71     return num_cleaned

```

```
72
73
74 #%% md
75 Do one simulation run with a simple randomized
    agent that has enough energy for 20 steps.
76 #%%
77 simple_environment(simple_randomized_agent,
    max_steps = 20)
78 #%% md
79 # Tasks
80
81 ## General [10 Points]
82
83 1. Make sure that you use the latest version of
    this notebook. Sync your forked repository and
    pull the latest revision.
84 2. Your implementation can use libraries like math
    , numpy, scipy, but not libraries that implement
    intelligent agents or complete search algorithms.
    Try to keep the code simple! In this course, we
    want to learn about the algorithms and we often do
    not need to use object-oriented design.
85 3. Your notebook needs to be formatted
    professionally.
86     - Add additional markdown blocks for your
        description, comments in the code, add tables and
        use matplotlib to produce charts where
        appropriate
87     - Do not show debugging output or include an
        excessive amount of output.
88     - Check that your PDF file is readable. For
        example, long lines are cut off in the PDF file.
        You don't have control over page breaks, so do not
        worry about these.
89 4. Document your code. Add a short discussion of
    how your implementation works and your design
    choices.
90
91
92 ## Task 1: Implement a simulation environment [20
    Points]
```

93  
94 The simple environment above is not very realistic.  
. Your environment simulator needs to follow the  
PEAS description from above. It needs to:  
95  
96 \* Initialize the environment by storing the state  
of each square (clean/dirty) and making some dirty  
. ([Help with random numbers and arrays in Python  
]([https://github.com/mhahsler/CS7320-AI/blob/master/Python\\_Code\\_Examples/random\\_numbers\\_and\\_arrays.ipynb](https://github.com/mhahsler/CS7320-AI/blob/master/Python_Code_Examples/random_numbers_and_arrays.ipynb)))  
97 \* Keep track of the agent's position.  
98 \* Call the agent function repeatedly and provide  
the agent function with the sensor inputs.  
99 \* React to the agent's actions. E.g., by removing  
dirt from a square or moving the agent around  
unless there is a wall in the way.  
100 \* Keep track of the performance measure. That is,  
track the agent's actions until all dirty squares  
are clean and count the number of actions it takes  
the agent to complete the task.  
101  
102 The easiest implementation for the environment is  
to hold an 2-dimensional array to represent if  
squares are clean or dirty and to call the agent  
function in a loop until all squares are clean or  
a predefined number of steps have been reached (i.  
e., the robot runs out of energy).  
103  
104 The simulation environment should be a function  
like the `simple\_environment()` and needs to work  
with the simple randomized agent program from  
above. \*\*Use the same environment for all your  
agent implementations in the tasks below.\*\*  
105  
106 \*Note on debugging:\* Debugging is difficult. Make  
sure your environment prints enough information  
when you use `verbose = True`. Also, implementing  
a function that the environment can use to  
displays the room with dirt and the current  
position of the robot at every step is very useful

```

106 .
107 #%%
108 import numpy as np
109 #print out the environment
110 def print_environment(environment):
111     print(np.matrix(environment))
112 #%%
113 # Your code and description goes here
114 def get_sensors(environment, botX, botY):
115     dimensions = len(environment)
116     sides = {"north" : False, "south" : False, "west" : False, "east" : False}
117     if botY == 0:
118         sides['west'] = True
119     if botX == 0:
120         sides['north'] = True
121     if botY == dimensions-1:
122         sides['east'] = True
123     if botX == dimensions-1:
124         sides['south'] = True
125     return sides
126
127 def check_clean(environment):
128     dimensions = len(environment)
129     for i in range(dimensions):
130         for j in range(dimensions):
131             if environment[i][j] == 'dirty':
132                 return False
133     return True
134
135
136 def environment(agent, steps, verbose = False,
137                 dimensions = 5):
138     #if steps is -1 it should run until the room
139     #is clean rather than the steps are reached
140     movements = {'north': {'x': -1, 'y': 0}, 'south': {'x': 1, 'y': 0}, 'east': {'x': 0, 'y': 1}, 'west': {'x': 0, 'y': -1}}
141     env = [ ['clean']*dimensions for i in range(dimensions)]
142     #assign dirt

```

```

141     for i in range(dimensions):
142         for j in range(dimensions):
143             if np.random.rand() < .2:
144                 env[i][j] = 'dirty'
145
146     #place vacuum
147     botX = np.random.randint(dimensions)
148     botY = np.random.randint(dimensions)
149     underTheBot = env[botX][botY]
150     env[botX][botY] = 'bot'
151
152     #do all steps for bot
153     if verbose:
154         print('initial')
155         print_environment(env)
156         print()
157
158     count = 0
159     if steps == -1:
160         while True:
161             walls = get_sensors(env, botX, botY)
162             dirt = False
163             if underTheBot == 'dirty':
164                 dirt = True
165             move = agent(walls, dirt)
166             if move == 'suck':
167                 underTheBot = 'clean'
168             else:
169                 if walls[move] is False:
170                     env[botX][botY] = underTheBot
171                     botX += movements[move]['x']
172                     botY += movements[move]['y']
173                     underTheBot = env[botX][botY]
174                     env[botX][botY] = 'bot'
175             if check_clean(env):
176                 if verbose:
177                     print('cleaned in', i, 'steps')
178             break
179             if verbose:
180                 print_environment(env)

```

```

181             print()
182             count += 1
183         return count
184
185     for i in range(steps):
186         walls = get_sensors(env, botX, botY)
187         dirt = False
188         if underTheBot == 'dirty':
189             dirt = True
190         move = agent(walls, dirt)
191         if move == 'suck':
192             underTheBot = 'clean'
193         else:
194             if walls[move] is False:
195                 env[botX][botY] = underTheBot
196                 botX += movements[move]['x']
197                 botY += movements[move]['y']
198                 underTheBot = env[botX][botY]
199                 env[botX][botY] = 'bot'
200             if check_clean(env):
201                 if verbose:
202                     print('cleaned in', i, 'steps')
203                     break
204                 if verbose:
205                     print_environment(env)
206                     print()
207                 count += 1
208
209             if verbose:
210                 print('after')
211                 print_environment(env)
212             return count, check_clean(env)
213
214
215 environment(simple_randomized_agent, 500)
216 %% md
217 ## Task 2: Implement a simple reflex agent [10 Points]
218
219 The simple reflex agent randomly walks around but reacts to the bumper sensor by not bumping into

```

```
219 the wall and to dirt with sucking. Implement the
agent program as a function.
220
221 _Note:_ Agents cannot directly use variable in the
environment. They only gets the percepts as the
arguments to the agent function.
222 #%%
223 # Your code and description goes here
224 def simple_reflex_agent(bumpers, dirty):
225     if dirty:
226         return 'suck'
227     open = []
228     for key, value in bumpers.items():
229         if value is False:
230             open.append(key)
231     return np.random.choice(open)
232
233 environment(simple_reflex_agent, 500)
234 #%% md
235 ## Task 3: Implement a model-based reflex agent [20 Points]
236
237 Model-based agents use a state to keep track of what they have done and perceived so far. Your agent needs to find out where it is located and then keep track of its current location. You also need a set of rules based on the state and the percepts to make sure that the agent will clean the whole room. For example, the agent can move to a corner to determine its location and then it can navigate through the whole room and clean dirty squares.
238
239 Describe how you define the __agent state__ and how your agent works before implementing it. ([Help with implementing state information on Python](# Your short description of the state and your
```

```
241 implementation goes here
242 # I used the kind we talked about by first
   navigating to the north-west corner and weaving up
   and down from there
243 # Global variables remember if I have started my
   weaving pattern or if I am still heading to the
   north-west corner and after that which direction I
   am currently weaving
244 #%%
245 # Your code goes here
246 started = False
247 heading = 'south'
248 # seen_corner = False
249
250 def model_based_agent(bumpers, dirty):
251     global started
252     global heading
253     #always clean if dirty
254     if dirty:
255         return 'suck'
256
257     #navigate towards the northwest corner
258     if not started and not bumpers['north']:
259         return 'north'
260
261     if not started and not bumpers['west']:
262         return 'west'
263
264     if not started and bumpers['north'] and
       bumpers['west']:
265         started = True
266         return 'south'
267
268     #start the weaving pattern
269     if bumpers[heading]:
270         heading = 'north' if heading == 'south'
271     else 'south'
272         return 'east'
273
274     else:
275         return heading
```

```
275 environment(model_based_agent, 500)
276 #%% md
277 ## Task 4: Simulation study [30 Points]
278
279 Compare the performance (the performance measure
    is defined in the PEAS description above) of the
    agents using environments of different size. E.g
    ., $5 \times 5$, $10 \times 10$ and
280 $100 \times 100$. Use 100 random runs for each.
    Present the results using tables and graphs.
    Discuss the differences between the agents.
281 ([Help with charts and tables in Python](https://
    github.com/mhahsler/CS7320-AI/blob/master/
    Python_Code_Examples/charts_and_tables.ipynb))
282 #%%
283 # Your code goes here
284 # I commented out the ones that would make this
    run all night... again
285 import matplotlib.pyplot as plt
286 random_five = 0
287 random_five_arr = []
288 random_ten = 0
289 random_ten_arr = []
290 random_hundred = 0
291 random_hundred_arr = []
292
293 reflex_five = 0
294 reflex_five_arr = []
295 reflex_ten = 0
296 reflex_ten_arr = []
297 reflex_hundred = 0
298 reflex_hundred_arr = []
299
300 model_five = 0
301 model_five_arr = []
302 model_ten = 0
303 model_ten_arr = []
304 model_hundred = 0
305 model_hundred_arr = []
306
307
```

```
308 for i in range(100):
309     # Add to random totals
310     tmp = environment(simple_randomized_agent, -1)
311     random_five += tmp
312     random_five_arr.append(tmp)
313     tmp = environment(simple_randomized_agent, -1,
314                         False, 10)
315     random_ten += tmp
316     random_ten_arr.append(tmp)
317     # tmp = environment(simple_randomized_agent, -1,
318     #                     False, 100)
319     # random_hundred += tmp
320     # random_hundred_arr.append(tmp)
321
322     # Add to reflex totals
323     tmp = environment(simple_reflex_agent, -1)
324     reflex_five += tmp
325     reflex_five_arr.append(tmp)
326     tmp = environment(simple_reflex_agent, -1,
327                         False, 10)
328     reflex_ten += tmp
329     reflex_ten_arr.append(tmp)
330
331     # Add to model totals
332     heading = 'south'
333     started = False
334     tmp = environment(model_based_agent, -1)
335     model_five += tmp
336     model_five_arr.append(tmp)
337
338     heading = 'south'
339     started = False
340     tmp = environment(model_based_agent, -1, False
341                         , 10)
342     model_ten += tmp
343     model_ten_arr.append(tmp)
```

```
344     heading = 'south'
345     started = False
346     tmp = environment(model_based_agent, -1, False
347     , 100)
347     model_hundred += tmp
348     model_hundred_arr.append(tmp)
349
350
351
352 random_five = int(random_five/100)
353 random_ten = int(random_ten/100)
354 random_hundred = int(random_hundred/100)
355
356 print('random five:', random_five)
357 print('random ten:', random_ten)
358 print('random hundred:', random_hundred)
359 print()
360
361 reflex_five = int(reflex_five/100)
362 reflex_ten = int(reflex_ten/100)
363 reflex_hundred = int(reflex_hundred/10)
364
365 print('reflex five:', reflex_five)
366 print('reflex ten:', reflex_ten)
367 print('reflex hundred:', reflex_hundred)
368 print()
369
370 model_five = int(model_five/100)
371 model_ten = int(model_ten/100)
372 model_hundred = int(model_hundred/100)
373
374 print('model five:', model_five)
375 print('model ten:', model_ten)
376 print('model hundred:', model_hundred)
377 print()
378
379 def show_graph(extra = ''):
380     x_pos = range(1,101)
381     global random_five_arr
382     global random_ten_arr
383     # global random_hundred_arr
```

```
384
385     global reflex_five_arr
386     global reflex_ten_arr
387     # global reflex_hundred_arr
388
389     global model_five_arr
390     global model_ten_arr
391     global model_hundred_arr
392
393     plt.plot(x_pos, random_five_arr, label = "
394         Random 5x5")
395     plt.plot(x_pos, random_ten_arr, label = "
396         Random 10x10")
397     if extra == 'random':
398         plt.plot(x_pos, random_hundred_arr, label
399             = "Random 100x100")
400
401     plt.plot(x_pos, reflex_five_arr, label = "
402         Reflex 5x5")
403     plt.plot(x_pos, reflex_ten_arr, label = "
404         Reflex 10x10")
405     if extra == 'reflex':
406         plt.plot(x_pos, reflex_hundred_arr, label
407             = "Reflex 100x100")
408
409     plt.xlabel('Iteration')
410     plt.ylabel('Steps Required')
411     plt.title('Steps Required Per Iteration By
412         Agent Type')
413
414     plt.legend()
415     plt.show()
```

```

415 show_graph()
416 #%% md
417 Fill out the following table with the average
    performance measure for 100 random runs (you may
    also create this table with code):
418
419 | Size      | Randomized Agent | Simple Reflex
    Agent | Model-based Reflex Agent |
420 |-----|-----|-----|
421 | 5x5      | 299           | 2518
                | 875493        |
422 | 10x10     | 112           | 982
                | 317621        |
423 | 100x100   | 25            | 120
                | 12096         |
424
425 Add charts to compare the performance of the
    different agents.
426 #%%
427 # Your graphs and discussion of the results goes
    here
428 # I have my graph output adjusted to omit the
    100x100 as they make the graph harder to read the
    other values because they are so much less
429 # However, the other 100x100s are fastest for the
    model, then the reflex, and extremely slow for
    the random agent. I ran the simulation overnight
    to get the number in the table above, but the
    number illustrate how much slower the random is
    than the reflex and the reflex is already
    significantly slower than the model
430 show_graph()
431 #I will also shoe the graph with the last model (
    100x100) since it still finishes in a reasonable
    amount of time compared to the reflex and random,
    it just makes it harder to visualize the results
432 show_graph('model')
433 #%% md
434 ## Task 5: Robustness of the agent implementations
    [10 Points]

```

```
435
436 Describe how **your agent implementations** will
   perform
437
438 * if it is put into a rectangular room with
   unknown size,
439 * if the cleaning area can have an irregular shape
   (e.g., a hallway connecting two rooms), or
440 * if the room contains obstacles (i.e., squares
   that it cannot pass through and trigger the bumper
   sensors).
441 #%%
442 # Answer goes here
443 # if it is put into a rectangular room with
   unknown size:
444 #      my bot will be able to clean this room in
   the expected amount of steps being (in an  $m \times n$ 
   room)  $mn * 1.2$  plus a little to get to the
   starting position for the model agent. This agent
   is consistent in this cleaning quantity and will
   always finish cleaning this room. The reflex is
   second fastest and does not have the guarantee of
   finishing; it also does not have a consistent
   quantity of solution steps. the attributes of the
   reflex bot also apply to the random bot, however
   it is far less efficient.
445
446 # if the cleaning area can have an irregular shape
   (e.g., a hallway connecting two rooms):
447 #      my model based agent would remain the
   fastest so long as it was allowed some
   modifications to understand the new possible
   movements. With a hallway, this would mean
   recognizing new spaces and routing through them
   accordingly. My two other agents, reflex and
   random, would see similar efficiency (so long as
   the bot had an equal chance at randomly happening
   into the hallway or new rooms as it did for just
   cleaning a single room, and this might not be the
   case with a narrow hallway the bot would have to
   find) and need no modification.
```

```
448
449 # if the room contains obstacles (i.e., squares
   that it cannot pass through and trigger the bumper
   sensors):
450 #      my reflex and random agents would require no
   modification. They might be slowed down if the
   obstacles block or make areas harder to reach and
   this could slow the efficiency of these already
   not very efficient bots would still, however, be
   able to finish the room. My model based
   implementation, however, would require some
   modification. if there was a corner shaped
   obstacle, it could trick the bot into thinking it
   had navigated to its starting position even if it
   hadn't yet. this would cause a premature start and
   the remaining left side of the room would stay
   dirty. in addition to this, the bot would
   inevitably hit some obstacles and think it had
   reached the top or bottom of the room and turn
   around, leaving the other side of the obstacle
   dirty. To solve this, I would need to implement
   some form of room mapping where the robot learns
   about where it is and if the obstacle was truly at
   the top or bottom of the room and if not, learn
   to go around it.
451 %% md
452 ## Graduate student advanced task: Obstacles [10
   Points]
453
454 __Undergraduate students:__ This is a bonus task
   you can attempt if you like [+5 Bonus Points].
455
456 1. Change your simulation environment to run
   experiments for the following problem: Add random
   obstacle squares that also trigger the bumper
   sensor. The agent does not know where the
   obstacles are. Observe how this changes the
   performance of the three implementations.
457
458 2. Describe what would need to be done to perform
   better with obstacles. Add code if you can.
```

```

459 #%%
460 # Your code and discussion goes here
461 def get_obstacle_sensors(environment, botX, botY):
462     dimensions = len(environment)
463     sides = {"north" : False, "south" : False, "
464         west" : False, "east" : False}
465         if botY == 0 or environment[botX][botY-1] == 'wall':
466             sides['west'] = True
467             if botX == 0 or environment[botX-1][botY] == 'wall':
468                 sides['north'] = True
469                 if botY == dimensions-1 or environment[botX][
470                     botY+1] == 'wall':
471                     sides['east'] = True
472                     if botX == dimensions-1 or environment[botX+1
473                         ][botY] == 'wall':
474                         sides['south'] = True
475             return sides
476 # I am making an environment that still has a .2
477 # chance of being dirty on some areas, but if it is
478 # not dirty, there is then a .2 chance a wall will
479 # be placed as an obstacle
480 # This env will return if it completed or not and
481 # if so how many steps it took to clean up
482 def obstacle_environment(agent, steps, verbose =
483     False, dimensions = 5):
484     movements = {'north': {'x': -1, 'y': 0}, ' '
485         'south': {'x': 1, 'y': 0}, 'east': {'x': 0, 'y': 1
486             }, 'west': {'x': 0, 'y': -1}}
487     env = [ ['clean']*dimensions for i in range(
488         dimensions)]
489     #assign dirt
490     for i in range(dimensions):
491         for j in range(dimensions):
492             if np.random.rand() < .2:
493                 env[i][j] = 'dirty'
494             elif np.random.rand() < .2:
495                 env[i][j] = 'wall'
496
497     #place vacuum

```

```
487     botX = np.random.randint(dimensions)
488     botY = np.random.randint(dimensions)
489     underTheBot = env[botX][botY]
490     env[botX][botY] = 'bot'
491
492     #do all steps for bot
493     if verbose:
494         print('initial')
495         print_environment(env)
496         print()
497
498     count = 0
499     for i in range(steps):
500         walls = get_obstacle_sensors(env, botX,
501                                         botY)
501         dirt = False
502         if underTheBot == 'dirty':
503             dirt = True
504         move = agent(walls, dirt)
505         if verbose:
506             print(walls)
507             print(move)
508         if move == 'suck':
509             underTheBot = 'clean'
510         else:
511             if walls[move] is False:
512                 env[botX][botY] = underTheBot
513                 botX += movements[move]['x']
514                 botY += movements[move]['y']
515                 underTheBot = env[botX][botY]
516                 env[botX][botY] = 'bot'
517             if check_clean(env):
518                 if verbose:
519                     print('cleaned in', i, 'steps')
520                     break
521             if verbose:
522                 print_environment(env)
523                 print()
524             count += 1
525
526     if verbose:
```

```
527     print('after')
528     print_environment(env)
529     return count, check_clean(env)
530
531 print('Below is the comparison of the old obstacle
      -less environment and the new, obstacle-ridden
      environment for 5x5 and 10x10 with a max 100,000
      steps (meaning it should finish unless not capable
      . Feel free to rerun and compare performances')
532 random_five_steps, random_five_done =
      obstacle_environment(simple_randomized_agent,
      100000)
533 random_ten_steps, random_ten_done =
      obstacle_environment(simple_randomized_agent,
      100000, False, 10)
534
535 reflex_five_steps, reflex_five_done =
      obstacle_environment(simple_reflex_agent, 100000)
536 reflex_ten_steps, reflex_ten_done =
      obstacle_environment(simple_reflex_agent, 100000,
      False, 10)
537
538 heading = 'south'
539 started = False
540 model_five_steps, model_five_done =
      obstacle_environment(model_based_agent, 100000)
541 heading = 'south'
542 started = False
543 model_ten_steps, model_ten_done =
      obstacle_environment(model_based_agent, 100000,
      False, 10)
544
545 print()
546 print('OBSTACLES: The random agent 5x5 finished in
      ' if random_five_done else 'OBSTACLES: The random
      agent 5x5 failed to finish in', random_five_steps
      , 'steps')
547 random_five_steps, random_five_done = environment(
      simple_randomized_agent, 100000)
548 print('The random agent 5x5 finished in' if
      random_five_done else 'The random agent 5x5 failed
```

```
548 to finish in', random_five_steps, 'steps')
549 print()
550 print('OBSTACLES: The random agent 10x10 finished
      in' if random_ten_done else 'OBSTACLES: The random
      agent 10x10 failed to finished in',
      random_ten_steps, 'steps')
551 random_ten_steps, random_ten_done = environment(
      simple_randomized_agent, 100000, False, 10)
552 print('The random agent 10x10 finished in' if
      random_ten_done else 'The random agent 10x10
      failed to finished in', random_ten_steps, 'steps')
553 print()
554 print('OBSTACLES: The reflex agent 5x5 finished in
      ' if reflex_five_done else 'OBSTACLES: The reflex
      agent 5x5 failed to finished in',
      reflex_five_steps, 'steps')
555 reflex_five_steps, reflex_five_done = environment(
      simple_reflex_agent, 100000)
556 print('The reflex agent 5x5 finished in' if
      reflex_five_done else 'The reflex agent 5x5 failed
      to finished in', reflex_five_steps, 'steps')
557 print()
558 print('OBSTACLES: The reflex agent 10x10 finished
      in' if reflex_ten_done else 'OBSTACLES: The reflex
      agent 10x10 failed to finished in',
      reflex_ten_steps, 'steps')
559 reflex_ten_steps, reflex_ten_done = environment(
      simple_reflex_agent, 100000, False, 10)
560 print('The reflex agent 10x10 finished in' if
      reflex_ten_done else 'The reflex agent 10x10
      failed to finished in', reflex_ten_steps, 'steps')
561 print()
562 print('OBSTACLES: The model agent 5x5 finished in'
      ' if model_five_done else 'OBSTACLES: The model
      agent 5x5 failed to finished in', model_five_steps
      , 'steps')
563 heading = 'south'
564 started = False
565 model_five_steps, model_five_done = environment(
      model_based_agent, 100000)
566 print('The model agent 5x5 finished in' if
```

```

566 model_five_done else 'The model agent 5x5 failed
      to finished in', model_five_steps, 'steps')
567 print()
568 print('OBSTACLES: The model agent 10x10 finished
      in' if model_ten_done else 'OBSTACLES: The model
      agent 10x10 failed to finished in',
      model_ten_steps, 'steps')
569 heading = 'south'
570 started = False
571 model_ten_steps, model_ten_done = environment(
      model_based_agent, 100000, False, 10)
572 print('The model agent 10x10 finished in' if
      model_ten_done else 'The model agent 10x10 failed
      to finished in', model_ten_steps, 'steps')
573 #%%
574 # It does not say I have to code this, section, so
      I will just say to guarantee the model solves it
      I would make sure it can avoid/prevent the
      following: if there was a corner shaped obstacle,
      it could trick the bot into thinking it had
      navigated to its starting position even if it hadn't
      yet. this would cause a premature start and the
      remaining left side of the room would stay dirty
      . in addition to this, the bot would inevitably
      hit some obstacles and think it had reached the
      top or bottom of the room and turn around, leaving
      the other side of the obstacle dirty. To solve
      this, I would need to implement some form of room
      mapping where the robot learns about where it is
      and if the obstacle was truly at the top or bottom
      of the room and if not, learn to go around it.
575
576 # The other two would not benefit from much
      different without changing the actual agent type
      from being the random or reflex
577 #%% md
578 ## More advanced implementation tasks
579
580 * __Agent for and environment with obstacles:__
      Implement an agent for an environment where the
      agent does not know how large the environment is (

```

```
580 we assume it is rectangular), where it starts or
where the obstacles are. An option would be to
always move to the closest unchecked/uncleaned
square (note that this is actually depth-first
search).
581
582 * __Utility-based agent:__ Change the environment
for a  $5 \times 5$  room, so each square has a
fixed probability of getting dirty again. For the
implementation, we give the environment a 2-
dimensional array of probabilities. The utility of
a state is defined as the number of currently
clean squares in the room. Implement a utility-
based agent that maximizes the expected utility
over one full charge which lasts for 100000 time
steps. To do this, the agent needs to learn the
probabilities with which different squares get
dirty again. This is very tricky!
583 #%%
584 # Your ideas/code
585
586 # I would run the depth first search (this could
be done using a camera or something on the actual
robot) then go to the nearest unclean square. if
there are no visible unclean squares, I would
ensure the entire room was visible to ensure 100%
clean this may not be the most efficient, but it
guarantees that nothing is needed to understand
current location and the robot will just always be
making meaningful motions towards a dirty square
587
588 # This seems really tricky, I think this would
require a large amount of robot training, so it
would get better at understanding when it would
need to turn around or come back later to a
particular square it decides has a high chance of
getting dirty again
```